

# 2017 Tubbs Fire Landscape Reconstruction

## 1 Validation objective and scope

This case asks whether ELMFIRE can reconstruct the landscape-scale spatial growth of the 2017 Tubbs Fire Landscape Reconstruction through 12 hours when driven by the case-local archived landscape, weather, moisture, ignition, spotting, and WUI inputs. The comparison evidence is a time-filtered set of VIIRS active-fire detections. The case exercises the coupled surface- and crown-fire, level-set, firebrand, ember-ignition, Monte Carlo, and building-spread pathways. It does not isolate or verify any single algorithm, and it does not establish predictive skill for events outside the three received reconstructions.

The scientific question is deliberately two-sided: does the ensemble overlap the observation-derived footprint, and does it do so without achieving overlap by excessive fire growth? Overlap, omission, commission, area, and probabilistic error are therefore reported together. The archive provides no independent, predeclared acceptance thresholds, so completed runs are *characterized*; they are not labeled PASS or FAIL.

## 2 Reference event, experiment, and observations

The configured event start is 8 October 2017 at 21:45 PDT, and the comparison instant is 9 October 2017 at 09:45 PDT (16:45 UTC). The observation directory contains 1168 point records in an EPSG:4326 shapefile; 597 occur on or before the comparison cutoff. The local fields identify the instrument as VIIRS and follow the NASA FIRMS active-fire schema. NASA describes VIIRS active-fire data as nominal 375-m thermal detections, not mapped fire perimeters. General product information is available from [https://firms.modaps.eosdis.nasa.gov/active\\_fire/](https://firms.modaps.eosdis.nasa.gov/active_fire/) and <https://www.earthdata.nasa.gov/data/catalog/lancemodis-vnp14img-nrt-2>.

The exact FIRMS collection, granules, download URL, acquisition date, filtering history, and license are not recorded in the case material. Consequently the NASA attribution is inferred from the field schema rather than established by a complete chain of custody. The local observation component checksums are recorded in `data/source_manifest.json`; no observation file is modified by preprocessing.

For raster comparison, detections through the cutoff are projected to the simulation CRS and enclosed by a concave hull with ratio 0.2 and no holes. If the installed geometry library lacks that operation, the postprocessor reports and uses a convex-hull fallback. This converts intermittent points into a comparison support; it is a declared workflow assumption, not an observed final perimeter. Cloud, smoke, overpass timing, geolocation error, false detections, and the hull choice remain observation uncertainties.




Figure 1: Case-local VIIRS detections retained through the comparison cutoff, colored by elapsed time from the configured event start.

### 3 Input data provenance and characterization

The received source is `Tubbs_VAL.rar` (SHA-256 `e97bf430d25122019211405f62f98979e929f80bc3f20d8c3b9f1f9f1a8a4955`). Its extracted payload is preserved under `data/archive_payload/`; every payload file is checked against `data/archive_manifest.json`. The collection’s only supplied provenance note is “From Dwi.” Upstream product names, providers, vintages, acquisition dates, processing history, and licenses for fuels, terrain, meteorology, moisture, WUI, and ignition data are unknown. The report therefore describes the archived values but does not infer unsupported product identities such as LANDFIRE or a specific weather analysis.

| Field          | Minimum | Mean  | Standard deviation | Maximum |
|----------------|---------|-------|--------------------|---------|
| Elevation      | 15      | 164.8 | 130.8              | 722     |
| Slope          | 0       | 8.329 | 8.383              | 47      |
| Wind speed     | 0.0608  | 24.32 | 13.19              | 54.81   |
| 1-h moisture   | 4.225   | 5.335 | 0.3679             | 6.793   |
| 10-h moisture  | 5.225   | 6.335 | 0.3679             | 7.793   |
| 100-h moisture | 6.225   | 7.335 | 0.3679             | 8.793   |

Table 1: Input statistics over valid cells and the 25 weather bands selected by the namelist. Wind speed is interpreted as 20-ft mph and dead-fuel moisture as percent from the archived flags and inspected current source; upstream metadata that would independently confirm these units is absent.

The 300-m weather rasters cover the same extent and CRS as the 30-m landscape but intentionally do not share its pixel grid. The preprocessor checks weather extent, band count, and CRS separately rather than treating the resolution difference as an error.

Aspect and wind direction are circular quantities. Their arithmetic means are not used as directional summaries; the circular means are  $351.5^\circ$  and  $40.21^\circ$ , respectively. Exact means and standard deviations use every valid selected value. Reported quantiles in `outputs/input_statistics.json` use a deterministic nearest-neighbor sample of at most  $512 \times 512$  cells per band, and are labeled accordingly.




Figure 2: Active elevation, slope, fuel-model, and ignition-mask inputs. These plots characterize model forcing; they are not validation evidence.




Figure 3: Distributions sampled deterministically from the 25 weather and dead-fuel-moisture bands selected by the namelist.




Figure 4: Exact domain-mean wind and dead-fuel-moisture values by active band.

## 4 Model configuration and rationale

The active domain is 700 by 900 cells at 30 m in the archived Lambert conformal conic CRS. Weather is 25-band, 300 m weather rasters at 30-minute intervals, and the ignition treatment is random ignition sampling from the archived ignition mask, with seed 2024. The simulation uses 50 members, a target CFL of 0.45, crown-fire model 2, and PER-MW firebrand generation with empirical transport distance, Eulerian accumulation, and physical ignition. These physical values, seeds, uncertainty ranges, ignition choices, and WUI settings are retained from the received archive except for the documented removal of a contradictory fixed Tubbs ignition. In the available material they are *documented by configuration* but their scientific selection rationale is unsupported; they must not be described as independently derived or literature-justified.

The archived namelist was migrated only where required by the inspected local ELMFIRE interface at base commit efa0b8fa3935ea4cd30a2bd054c7e9b83db81163, whose working tree already contained modifications. The obsolete computational-domain group and DEBUG\_LEVEL entry were removed, and HRR\_ELLIPSE\_ADJ was moved to WUI. Legacy spotting booleans were mapped to PER-MW generation, empirical distance, Eulerian accumulation, and physical ignition. Archived building spread type 3 was mapped to current UCB/UMD type 2. The contradictory archived fixed ignition was disabled because current ELMFIRE applies it in addition to the configured random ignition-mask

sample, whereas this case specifies the latter alone. The complete mapping and audit limitations are recorded in the suite-level `SOURCE_AUDIT.md`. The active assignments pass the repository namelist schema and the test-critical values are locked in `case.yaml`.

| Choice                                                              | Evidence class    | Basis and limitation                                                                                                                                |
|---------------------------------------------------------------------|-------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------|
| Landscape, weather, ignition, Monte Carlo, spotting, and WUI values | documented        | Retained from the received archive; upstream derivation and event-specific tuning history are unknown.                                              |
| Namelist interface mappings                                         | derived           | Required by identifiers and enumerations in the inspected current source; no physical values were tuned.                                            |
| Comparison time                                                     | documented        | Locked in the case contract and consistent with the reconstructed archive workflow; original scientific rationale is absent.                        |
| $p \geq 0.50$ footprint                                             | workflow-selected | Majority-ensemble interpretation used for reproducibility; it is not an acceptance threshold and sensitivity should be reported for scientific use. |
| VIIRS hull ratio 0.2                                                | workflow-selected | Compact cumulative support for point detections; not calibrated to an uncertainty-bounded perimeter.                                                |

Table 2: Provenance of consequential configuration and comparison choices.

## 5 Preprocessing and reproducibility

The case-local `scripts/preprocess.py` verifies the source archive, every materialized payload file, and every observation component by size and SHA-256. It checks that required landscape rasters share a grid, confirms that the namelist-selected weather rasters contain at least 25 bands, computes input and observation statistics, and regenerates the four input figures. It writes only `outputs/input_statistics.json` and `figures/*.pdf`; it does not rewrite raw evidence.

`run_case.sh` then launches ELMFIRE, runs the case-local postprocessor, and compiles this report. The executable, Python interpreter, and MPI rank count can be selected through `ELMFIRE_BIN`, `PYTHON_BIN`, and `ELMFIRE_MPI_RANKS`. The default 50 ranks assign one rank to each ensemble member and avoid an unnecessary virtual rank. Outputs and logs are retained rather than deleted or renamed. Reproducible input identity does not by itself guarantee result reproducibility: each run must also record the actual ELMFIRE revision, compiler, dependencies, and platform.

## 6 Validation metrics and comparison framework

For member  $i$ , let  $\tau_i(\mathbf{x})$  be the time-of-arrival raster and  $t_c$  the declared comparison time. The ensemble burn probability and thresholded footprint are

$$p(\mathbf{x}) = \frac{1}{N} \sum_{i=1}^N \mathbf{1}[0 < \tau_i(\mathbf{x}) \leq t_c], \quad P = \{\mathbf{x} : p(\mathbf{x}) \geq 0.50\}. \quad (1)$$

For the rasterized observation support  $O$ , intersection  $I = |P \cap O|$ , and valid-cell count  $M$ , the workflow reports

$$J = \frac{|P \cap O|}{|P \cup O|}, \quad \text{precision} = \frac{I}{|P|}, \quad \text{recall} = \frac{I}{|O|}, \quad (2)$$

$$F_1 = \frac{2 \text{ precision recall}}{\text{precision} + \text{recall}}, \quad R_A = \frac{|P|}{|O|}, \quad \text{Brier} = \frac{1}{M} \sum_x (p - y)^2. \quad (3)$$

Jaccard is a symmetric footprint-overlap summary. Precision diagnoses commission/overgrowth, recall diagnoses omission/undergrowth, and  $F_1$  balances those errors. The area ratio exposes size compensation that overlap alone can hide. The Brier score retains ensemble probabilities, although the large unburned domain can dominate it. Burned-area growth and conditional median arrival time add temporal context. VIIRS detections and their derived hull are the observational baseline; no previous simulation result is used as truth.

Metrics are calculated only when exactly 50 uniquely identified members reach the comparison time and share one grid and CRS. This completeness gate prevents partial runs from being reported as scientific results. The static aspect raster defines the analysis-domain mask. Member-specific NoData cells in time-of-arrival outputs represent cells not burned by that member; they are therefore retained as zero burn outcomes rather than treated as a grid-mask mismatch.

## 7 Simulation results and field-level comparison

The current workflow status is **CHARACTERIZED**. Metrics are descriptive because no independent predeclared acceptance thresholds are available. Metrics and figures are generated only from the complete, current case-local ensemble. If that ensemble is unavailable, the workflow reports the result as not evaluable rather than substituting legacy outputs or zero-valued metrics.




Figure 5: Ensemble burn probability, VIIRS-derived support, and cell-level true-positive, false-positive, and false-negative regions.

2017 Tubbs Fire landscape reconstruction: arrival time where at least one member burns




Figure 6: Conditional median arrival time among members that burn each cell by the comparison time. Conditioning is stated because it omits no-burn members.

2017 Tubbs Fire landscape reconstruction: available final-time fields (representative member)




Figure 7: Applicable final-time two-dimensional fields found by the postprocessor. Omission of this panel means no recognized optional field was available; it does not imply a zero field.

The active deck requests flame length and ember flux in addition to arrival time. It does not request a structure-damage count or damage-time output, so a damaged-structures time series is not reported

and must not be inferred from WUI input rasters alone.

## 8 Integrated statistics and calculated metrics

| Metric                                      | Current value |
|---------------------------------------------|---------------|
| Jaccard overlap                             | 0.555         |
| Precision                                   | 0.6051        |
| Recall                                      | 0.8702        |
| $F_1$ score                                 | 0.7138        |
| Brier score                                 | 0.09008       |
| Predicted/observed area ratio               | 1.438         |
| Predicted area (km <sup>2</sup> )           | 133.8         |
| Observation-support area (km <sup>2</sup> ) | 93.02         |
| Ensemble members found                      | 50            |

Table 3: Calculated statistics. **MISSING** is intentionally conspicuous when the completeness gate prevents evaluation.




Figure 8: Median and 5th–95th percentile ensemble burned-area growth.

These values are compared with the observation-derived support and not with a stored ELMFIRE baseline. Because neither metric thresholds nor observation uncertainty bounds were declared independently, a populated table supports scientific characterization only.

## 9 Interpretation and uncertainty

A low Jaccard value with high recall and low precision indicates spatial overgrowth; high precision with low recall indicates undergrowth or a displaced trajectory. An area ratio near one cannot establish spatial agreement. The Brier score must be read with the footprint metrics because correct prediction of the large unburned domain may obscure errors near the fire.

The dominant limitations are incomplete input provenance, an observation product that samples thermal anomalies rather than perimeter, uncertain hull sensitivity, unknown event-specific calibration history, and absence of predeclared acceptance bounds. Model-form error, uncertain fuels and moisture, meteorological representation, ignition timing/location, spotting parameters, and WUI parameters are not separable using these aggregate metrics alone. The 50 archived perturbations characterize one configured ensemble; they do not by themselves quantify all epistemic uncertainty.

## 10 Conclusions and limitations

The case is a self-contained, auditable validation workflow: raw archive and observation identities are checked, inputs are statistically characterized, output completeness is enforced, applicable raster and integrated results are visualized, and metric definitions and limitations are explicit. Its current scientific result is **CHARACTERIZED**; the accompanying reason states whether the complete current ensemble supported quantitative characterization.

To support a defensible validation decision, future case owners must supply the upstream input citations and licenses, exact VIIRS collection/granule and acquisition provenance, the derivation or literature basis for ignition, spotting, WUI, and perturbation settings, disclosure of any calibration against this event, and acceptance thresholds fixed before inspecting the evaluated run. Until then, even a completed run is correctly labeled **CHARACTERIZED**, not **PASS** or **FAIL**.